



BIO306: Bioinformatics

Lecture 2


NGS and Reads mapping

Wenfei JIN PhD

jinwf@sustc.edu.cn

Department of Biology, SUSTech

Sanger Sequencing



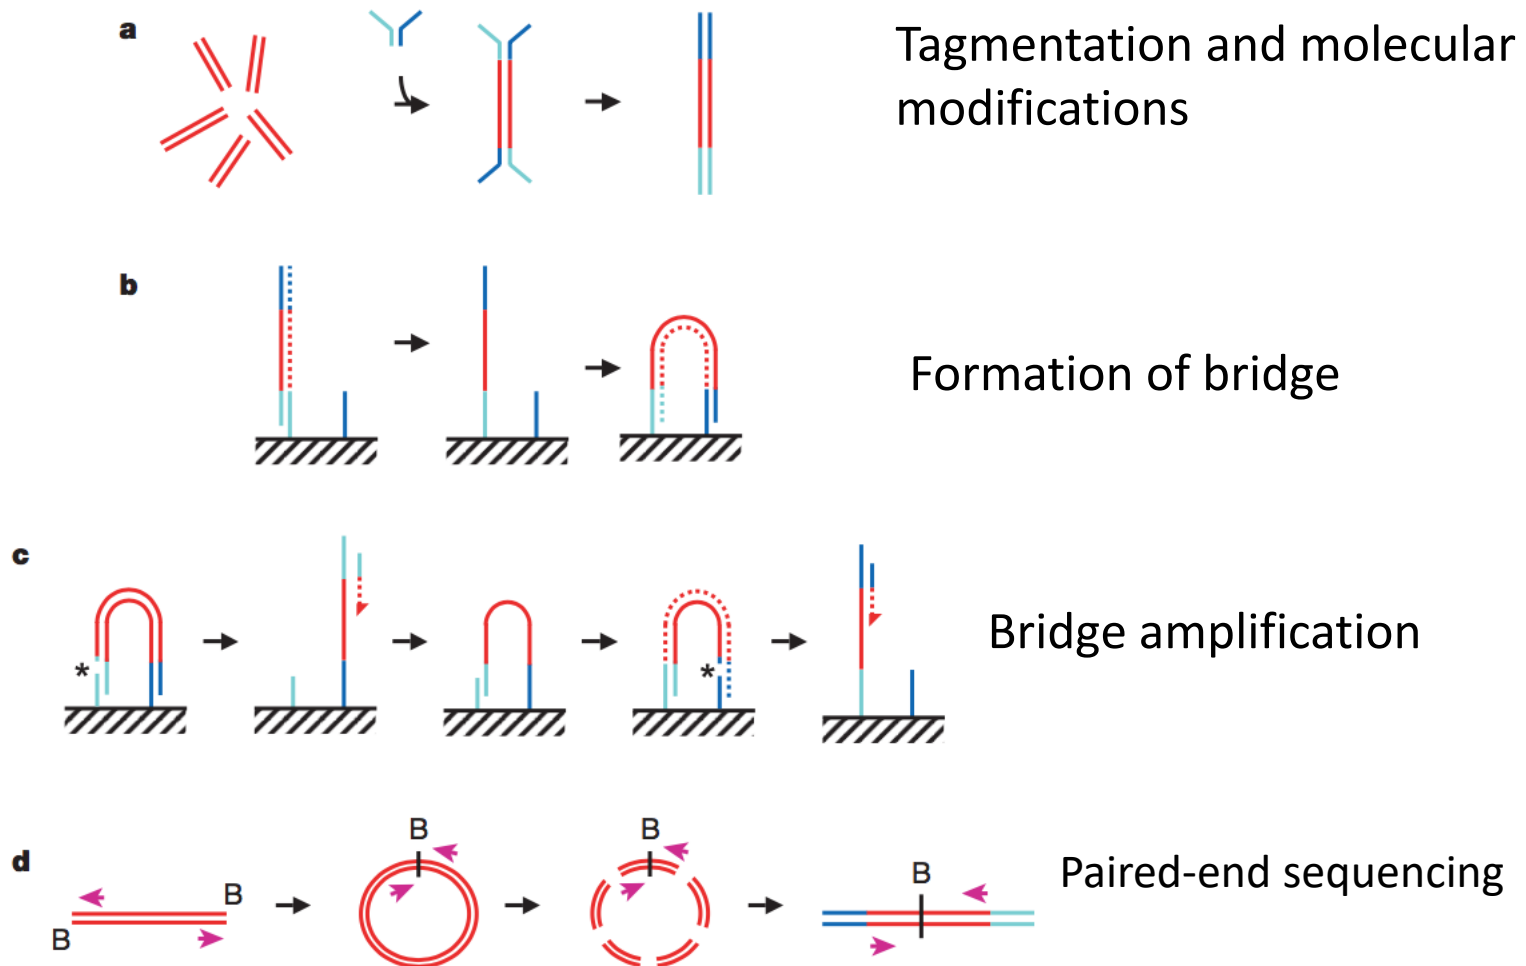
Progression of Sequencing Reaction

dideoxynucleotides (ddNTPs)




What is Next generation sequencing (NGS)?

- High-throughput sequencing
- Massively parallel sequencing
- Illumina dye sequencing as example



Illumina sequencing showed in original nature paper



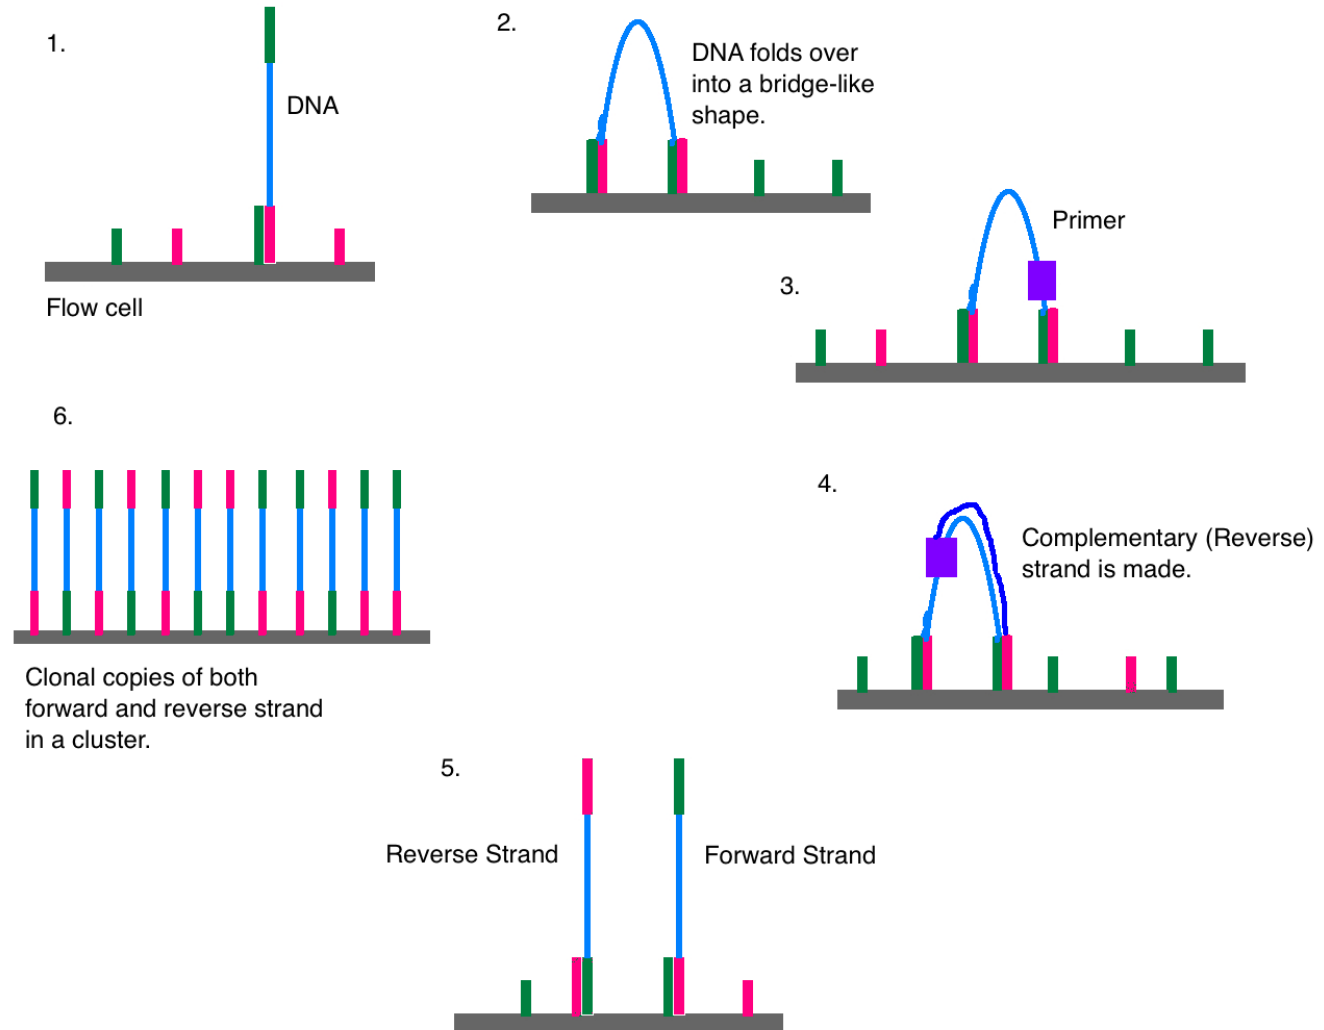
Flow cell




Library preparation



Amplification/Cluster generation

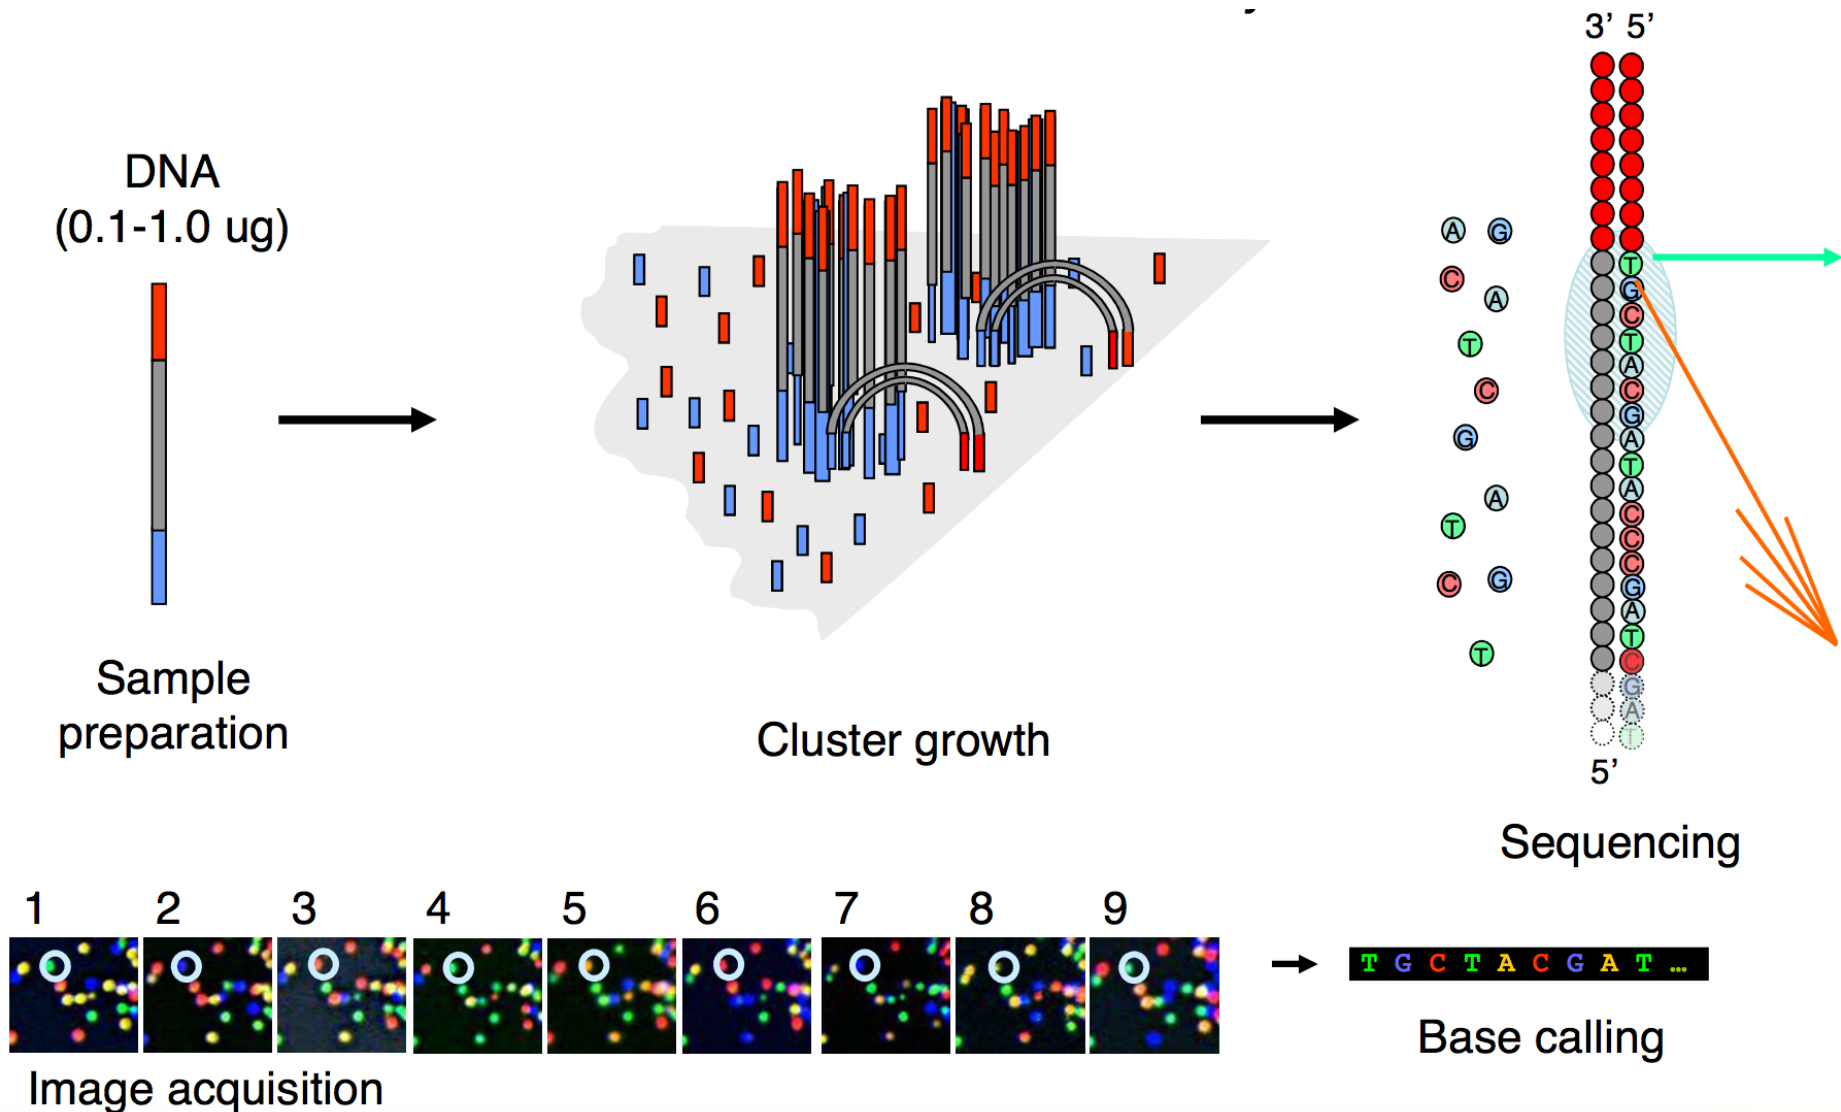


Sequence by Synthesis (1)

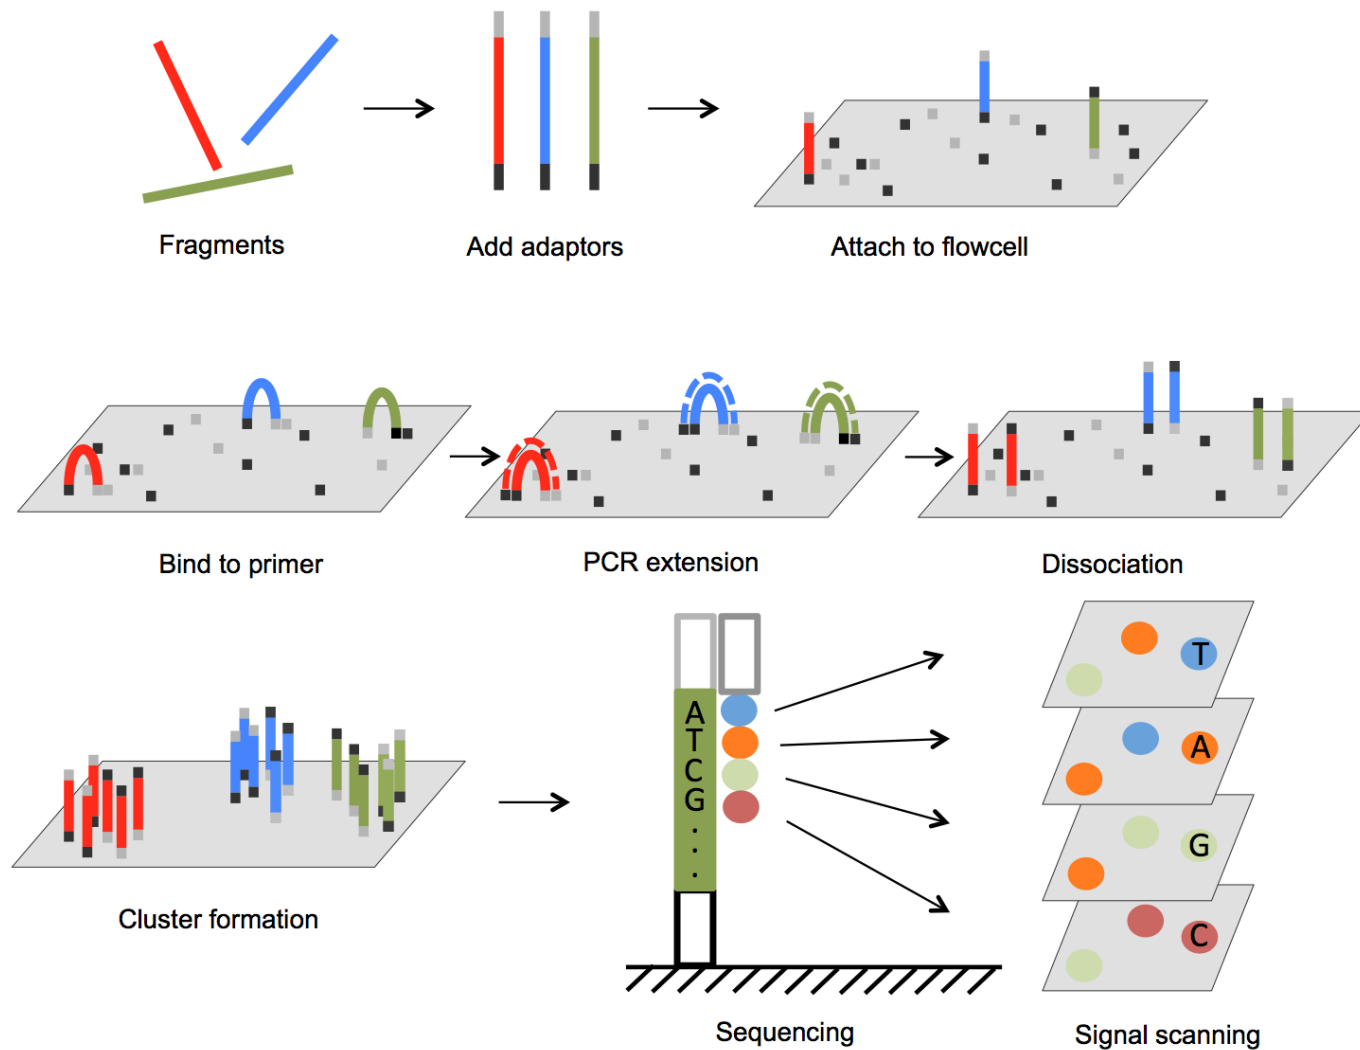


fluorescently tagged nucleotides to the DNA strand

Sequence by Synthesis (2)



Schematic of Illumina dye sequencing



Illumina sequencing

DNA Sample



Construct
Library



Cluster Generation
in Flow Cell



Sequencing by
Synthesis



200+ million reads per lane
(>100 bp reads)

Data Processing

Nucleotide Flows



Raw Images



Image Processing

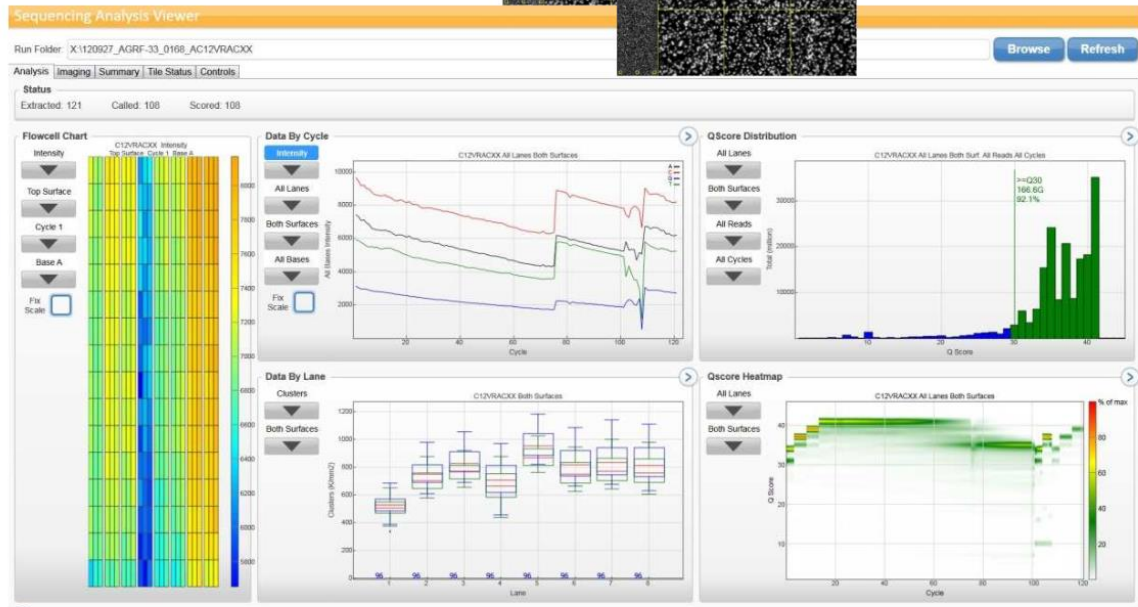

Base-calling

Quality Filtering

.bcl



Fastq



Fastq format

@SEQ_ID

GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAA

+

!"*(((('*'+))%%%++)(%%%%).1***+*"))**55CCF>>>>>>

4 lines per sequence/read

Line 1 begins with a @ and is followed by ID

Line 2 sequence letters.

Line 3 begins with a '+' is optionally followed by any characters

Line 4 quality values for the sequence in Line 2,

must contain the same number of symbols as letters in the sequence.

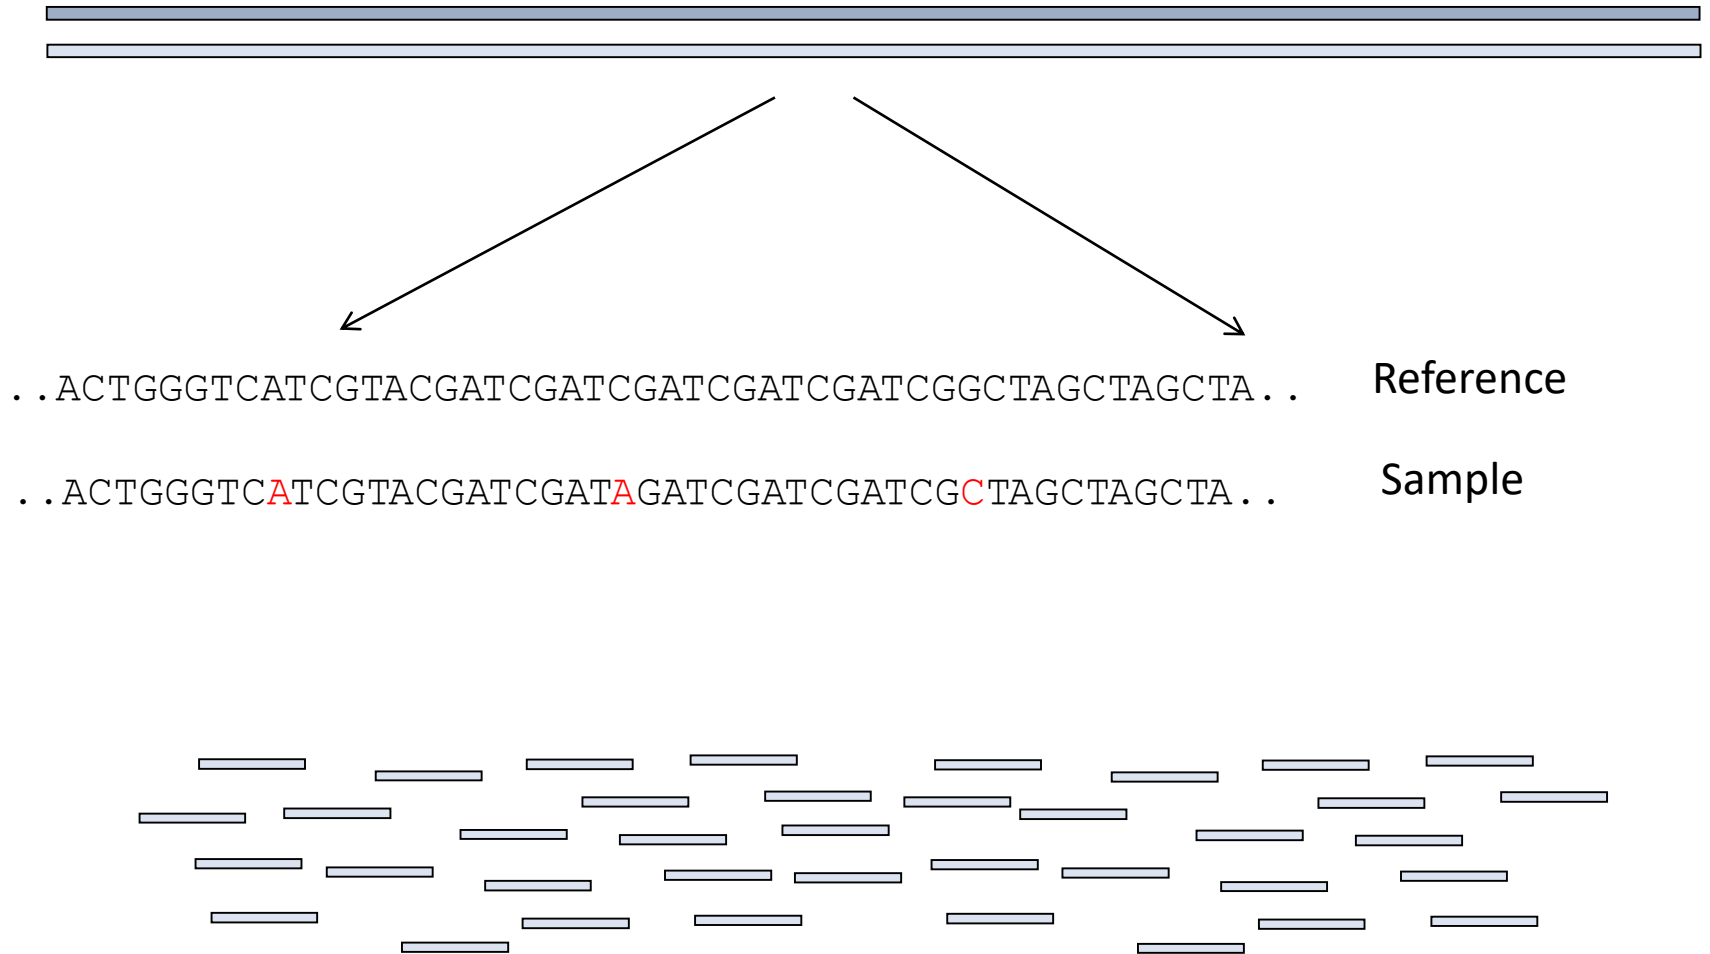
quality score are each encoded with a single ASCII character for brevity.

Characters of NGS data

- Short reads
 - Illumina (36 – 300bp)
 - SoLID (75bp max)
 - Ion Torrent (200-300bp max – currently...)
 - Roche 454 – 400-800bp
- Data set is large (multiple coverage)
 - Millions or even billions reads

Reads alignment

Alignment of reads to a reference



Short Read Applications

- Genotyping

Goal: identify variations

...CCATAG TATGCGCCC CGGAATT GGTATAC...
...CCAT CTATATGCG TCGGAATT CGGTATAC
...CCAT GGCTATATG CTATCGGAAA CGGTATAC
...CCA AGGCTATAT CCTATCGGA GCGGTATA C...
...CCA AGGCTATAT GCCCTATCG TTTGCGGT C...
...CC AGGCTATAT GCCCTATCG AAATTTGC ATAC...
...CC TAGGCTATA GCGCCCTA AAATTTGC GTATAC...

...CCATAGGCTATATGCGCCCTATCGGCAATTTGCGGTATAC...

- RNA-seq, ChIP-seq, Methyl-seq

Goal: classify, measure significant peaks

GAAATTTGC
GGAAATTTG
CGGAAATTT
CGGAAATTT
TCGGAAATT
CTATCGGAAA
CCTATCGGA TTTGCGGT
GCCCTATCG AAATTTGC
GCCCTATCG AAATTTGC ATAC...

...CC

...CCATAGGCTATATGCGCCCTATCGGCAATTTGCGGTATAC...

Why is short read alignment hard?

The shorter a read, the less likely it is to have a unique match to a reference sequence




Fig. 1 The proportion of unique sequence in the *Streptococcus suis* (squares) and *Mus musculus* (triangles) genomes for varying read lengths. This graph indicates that read length has a critical affect on the ability to place reads uniquely to the genome

Why do we generate short reads?

- Sanger reads lengths ~ 800-2000bp
- Generally we define short reads as anything below 200bp
 - Illumina (100bp – 250bp)
 - SOLiD (75bp max)
 - Ion Torrent (200-300bp max – currently...)
 - Roche 454 – 400-800bp
- Even with these platforms it is cheaper to produce short reads (e.g. 50bp) rather than 100 or 200bp reads
- Diminishing returns:
 - For some applications 50bp is more than sufficient
 - Resequencing of smaller organisms
 - Bacterial de-novo assembly
 - ChIP-Seq
 - Digital Gene Expression profiling
 - Bacterial RNA-seq

Contents

- **Alignment algorithms for short-reads**
 - Adapting hashed seed-extend algorithms to work with shorter reads
 - Indel detection
 - Suffix/Prefix Tries
 - Other alignment considerations
 - Typical alignment pipeline
- **Assembly algorithms for short reads**
 - Effect of repeats
 - Overlap-Consensus
 - de Bruijn graphs
 - Assembly evaluation metrics
 - Typical assembly pipeline

Adapting hashed seed-extend algorithms to work with shorter reads

- Improve seed matching sensitivity
 - Allow mismatches within seed
 - BLAST
 - Allow mismatches + Adopt spaced-seed approach
 - ELAND, SOAP, MAQ, RMAP, ZOOM
 - Allow mismatches + Spaced-seeds + Multi-seeds
 - SSAHA2, BLAT, ELAND2
- Above and/or Improve speed of local alignment for seed extension
 - Single Instruction Multiple Data
 - Shrimp2, CLCBio
 - Reduce search space to region around seed

Hashed seed-extend algorithms

- **2 step process**
 - Identify a match to the seed sequence in the reference
 - Extend match using sensitive (but slow) Smith-Waterman algorithm (dynamic programming)

Seed-extend algorithm

Reference sequence:

. . . ACTGGGTCATCGTACGATCGATCGATCGATCGATCGGCTAGCTAGCTA . . .

Short read:

GTCATCGTACGATCGATAGATCGATCGGCTA

Note that the short read has 1 difference wrt to reference

Seed-extend algorithm

Reference sequence:

...ACTGGGTCATCGTACGATCGATCGATCGATCGGCTAGCTAGCTA...

Short read:

GTCATCGTACG ATCGATAGATCG ATCGATCGGCTA

11bp word

11bp word

11bp word

The algorithm will try to match each word to the reference. If there is a match at with any single word it will perform a local alignment to extend the match

Seed-extend algorithm

Reference sequence:

Seed Extend with Smith Waterman

...ACTGGGTCATCGTACGATCGATCGATCGATCGATCGGCTAGCTAGCTA...

GTCATCGTACGATCGAACGATCGATCGATCGGCTA

Short read:

GTCATCGTACG ATCGATAGATCG ATCGATCGGCTA

Here the algorithm is able to match the short read with a word length of 11bp

Seed-extend algorithm

Reference sequence:

. . . ACTGGGTCATCGTACGATCGATCGATCGATCGATCGGCTAGCTAGCTA . . .

Short read:

GTCATC GTACGATCGATCGATCGGCA

Note that the short read has 3 differences
Possibly sequencing errors, possibly SNPs

Seed-extend algorithm

Reference sequence:

. . . ACTGGGTCATCGTACGATCGATCGATCGATCGATCGGCTAGCTAGCTA . . .

Short read:

GTCATC GTACG

11bp word

ATCGATC GATCG

11bp word

ATCGATCGGC AA

11bp word

Note that the short read has 3 differences

Seed-extend algorithm

Reference sequence:

. . . ACTGGGTCATCGTACGATCGATCGATCGATCGATCGGCTAGCTAGCTA . . .

Short read:

GTCAT^CGTACG ATCGATC^GATCG ATCGATCGGC^AA

No seeds match

Therefore the algorithm would find no hits at all!

Adapting hashed seed-extend algorithms to work with shorter reads

- Improve seed matching sensitivity
 - **Allow mismatches within seed**
 - **BLAST**
 - Allow mismatches + Adopt spaced-seed approach
 - ELAND, SOAP, MAQ, RMAP, ZOOM
 - Allow mismatches + Spaced-seeds + Multi-seeds
 - SSAHA2, BLAT, ELAND2
- Above and/or Improve speed of local alignment for seed extension
 - Single Instruction Multiple Data
 - Shrimp2, CLCBio
 - Reduce search space to region around seed

Adapting hashed seed-extend algorithms to work with shorter reads

- Improve seed matching sensitivity
 - **Allow mismatches within seed**
 - **BLAST**
 - **Allow mismatches + Adopt spaced-seed approach**
 - **ELAND, SOAP, MAQ, RMAP, ZOOM**
 - Allow mismatches + Spaced-seeds + Multi-seeds
 - SSAHA2, BLAT, ELAND2
- Above and/or Improve speed of local alignment for seed extension
 - Single Instruction Multiple Data
 - Shrimp2, CLCBio
 - Reduce search space to region around seed

Contents


- **Alignment algorithms for short-reads**
 - Background – Blast (why can't we use it?)
 - Adapting hashed seed-extend algorithms to work with shorter reads
 - Suffix/Prefix Tries**
 - Other alignment considerations
 - Typical alignment pipeline
- **Assembly algorithms for short reads**
 - Effect of repeats
 - Overlap-Consensus
 - de Bruijn graphs
 - Assembly evaluation metrics
 - Typical assembly pipeline

Suffix-Prefix Trie

- A family of methods which uses a Trie structure to search a reference sequence
 - Bowtie
 - BWA
 - SOAP version 2
- Trie – data structure which stores the suffixes (i.e. ends of a sequence)
- Key advantage over hashed algorithms:
 - Alignment of multiple copies of an identical sequence in the reference only needs to be done once
 - Use of an FM-Index to store Trie can drastically reduce memory requirements (e.g. Human genome can be stored in 2Gb of RAM)
 - Burrows Wheeler Transform to perform fast lookups

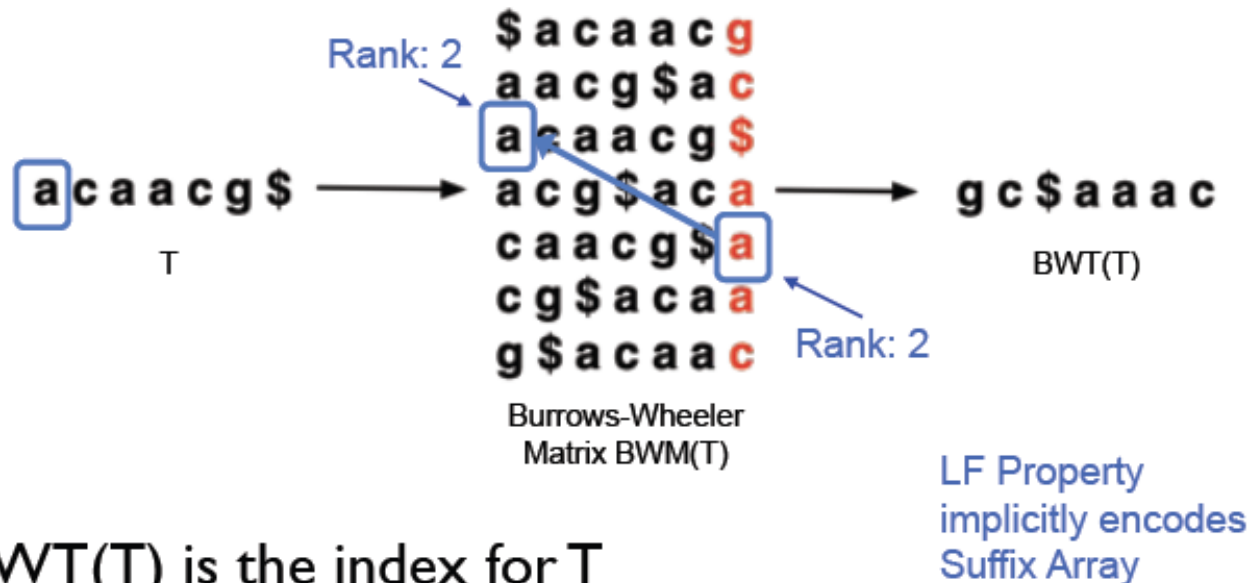
Suffix Trie

AGGAGC



Heng Li & Nils Homer.
Sequence alignment
algorithms for next-
generation sequencing.
Briefings in
Bioinformatics. Vol 11.
No 5. 473 483, 2010

Suffix Trie



- $BWT(T)$ is the index for T

A block sorting lossless data compression algorithm.

Burrows M, Wheeler DJ (1994) *Digital Equipment Corporation*. Technical Report 124

Burrows-Wheeler Algorithm

- Encodes data so that it is easier to compress
- Burrows-Wheeler transform of the word BANANA
- Can later be reversed to recover the original word

Transformation				
Input	All Rotations	Sorting All Rows in Alphabetical Order by their first letters	Taking Last Column	Output Last Column
<div>^BANANA </div>	<div>^BANANA </div> <div> ^BANANA</div> <div>A ^BANAN</div> <div>NA ^BANA</div> <div>ANA ^BAN</div> <div>NANA ^BA</div> <div>ANANA ^B</div> <div>BANANA ^</div>	<div>ANANA ^B</div> <div>ANA ^BAN</div> <div>A ^BANAN</div> <div>BANANA ^</div> <div>NANA ^BA</div> <div>NA ^BANA</div> <div>^BANANA </div> <div> ^BANANA</div>	<div>ANANA ^B</div> <div>ANA ^BAN</div> <div>A ^BANAN</div> <div>BANANA ^</div> <div>NANA ^BA</div> <div>NA ^BANA</div> <div>^BANANA </div> <div> ^BANANA</div>	<div>BNN^AA A</div>

More Burrows-Wheeler

Input

SIX.MIXED.PIXIES.SIFT.SIXTY.PIXIE.DUST.BOXES

Burrows-Wheeler Output

TEXYDST.E.IXIXIXSSMPPS.B..E.S.EUSFXDIIIOIIIT

Repeated characters mean that it is easier to compress

Bowtie/Soap2 example

Reference



BWT(Reference)

Query:

AATGATACGGCGACCACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATGATACGGCGACCACCGAGATCTA



Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATGATACGGCGACCACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATGATACGGCGACCACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)

Query:

AATGATACGGCGAC **CACCGAGATCTA**

Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATGATACGGCGACCACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)

Query:

AATGATACGGCGACCAACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATG TACGGCGACCAACCGAGATCTA

Bowtie/Soap2 example

Reference



BWT(Reference)



Query:

AATGTTACGGCGACCACCGAGATCTA

Bowtie/Soap2 vs. BWA

- Bowtie and Soap2 cannot handle gapped alignments
 - No indel detection => Many false SNP calls

Bowtie/Soap2:

ACTCCCATTGTCATCGTACTTGGGATC^GTAACA Reference

CCATTGTCATCGTACTTGGGATC^{TA}

TCATCGTACTTGGGATC^{TA}

TTGGGATC^{TA}

↖ False SNPs

N.B. Bowtie2 can handle gapped alignments

Bowtie/Soap2 vs. BWA

- Bowtie and Soap2 cannot handle gapped alignments
 - No indel detection => Many false SNP calls

BWA:

ACTCCCATTGTCATCGTACTTGGGATC~~G~~TAACA Reference

CCATTGTCATCGTACTTGGGATC~~TA~~

TCATCGTACTTGGGATC~~TA~~

TTGGGATC~~TA~~

N.B. Bowtie2 can handle gapped alignments

Comparison

Hash referenced spaced seeds

- Requires ~50Gb of memory
- Runs 30-fold slower
- Is much simpler to program
- Most sensitive

Suffix/Prefix Trie

- Requires <2Gb of memory
- Runs 30-fold faster
- Is much more complicated to program
- Least sensitive

Comparison

- Bowtie's reported 30-fold speed increase over hash-based MAQ with small loss in sensitivity
- Limitations to Trie-based approaches:
 - Only able to find alignments within a certain 'edit distance'
 - Bowtie does not do gapped alignments – no indels!
 - Important to quality clip reads (-q in BWA)
 - Non-A/C/G/T bases on reads are simply treated as mismatches
 - Make sure Ns are removed!

Hash based approaches are more suitable for divergent alignments

- Rule of thumb:
 - <2% divergence -> Trie-based
 - E.g. human alignments
 - >2% divergence -> seed-extend based approach
 - E.g. wild mouse strains alignments

Thank you for your attention!